STAKEHOLDER ONE-PAGER

A Governed AI Assistant for Engineering

What it saves, what it costs and what could go wrong, in plain terms.

Vijayalakshmi Ramu | October 2026 | github.com/Viji-qmofi/first-microservices-Viji-R-agentic-harness-cicd

45+ → 2.5 min
Human time per build check: by hand (estimate) vs reviewing the AI's report (timed once)

$0.38–$0.81
AI cost per build check (measured)

264× faster
One step turned into plain code; its cost fell from $0.17 to $0

4 of 4
Deliberate attempts to break the rules were blocked

THE PROBLEM

Before a team ships a change to its four-service application, someone runs the builds, reads the logs and decides whether it is safe to go on. By hand that takes at least 45 minutes (the reviewer's own estimate). It is repetitive, and tired readers miss warnings.

WHAT IT IS

A small team of AI roles does the first pass: one plans, one builds, two review, and a manager role checks each result instead of trusting it. One step, reading the test results, is ordinary code, not AI. **A person approves before anything is saved.** Each role can use only the tools its job needs, and the system enforces that, not just a written policy.

RISKS AND HOW THEY ARE MANAGED

- **A wrong answer accepted.** A person reviews every report. A scored check against a known answer caught a false "no warnings" report.
- **An AI role overreaching.** Attempts to act outside a role were blocked in 4 of 4 tests. A write that once got through is now denied.
- **Runaway cost or time.** Limits of $2.00 and 15 minutes per task, checked after every run. The build pipeline cancels any job past its time limit (shown on a test run).
- **Secrets.** Keys are never stored in the repository, and a commit check blocks credential-like text.

WHAT IT IS WORTH

Assumptions, not measurements: a developer costs $75 an hour and the team runs 20 checks a week.

- **If every report were right:** about 42 minutes saved per check, roughly **14 hours and $1,060 a week**, less about $10 of AI cost.
- **At today's measured reliability** (1 report in 3 passes): about 12 minutes saved per check, roughly **4 hours and $310 a week**. This assumes every bad report is caught and redone by hand.
- **The price of control:** governance adds about 32% to the AI cost (about 12 cents) and about a minute per run.

HONEST LIMITS

- Only **1 in 3** reports passed on this task with today's wording. The fix is known (replace the log-reading step with a fixed rule) and is not built yet.
- The hours and dollars are projections from an estimated manual baseline, not measured on a live team.
- Tested on a representative application, not customer data. The planning and building roles are written for one service, the order service. Role checks are policy against declared roles, not proof of identity.

NEXT STEPS

1. Replace the log-reading step with a fixed rule, then re-measure.
2. Time reviews by a second person who has not seen the answer key.
3. Pilot on a live team's tickets, with the approvals real data requires.